

Supplementary Material

Taylor-Made Sparsity

Anonymous Authors · NeurIPS 2026 Submission

This document provides supplementary material to accompany the main paper. It contains qualitative Grad-CAM visualizations (Appendix A), extended raw data tables (Appendix B), and full experimental setup details including hardware specifications and hyperparameters (Appendix C).

Appendix A

Qualitative Interpretability Visualizations

Each row below shows a Grad-CAM heatmap comparison across four model variants for a single CIFAR-10 test image: the unpruned **Baseline**, **L1-Norm** (magnitude pruning), **Grad-Act (Ours)** (Taylor-Made Sparsity), and **DepGraph** (dependency-graph baseline). All pruned models use a 15% filter removal ratio. We observe that Grad-Act preserves localized high-response regions corresponding to object boundaries, whereas L1-based pruning disperses activation mass across background regions, indicating loss of discriminative features. This visually corroborates our central claim: the Taylor proxy selects filters that are geometrically meaningful to the task, not merely those with large weight magnitudes.




Figure A.1. Sample #000 — Baseline · L1-Norm · Grad-Act (Ours) · DepGraph at 15% sparsity. Warmer colours indicate higher activation.




Figure A.2. Sample #002 — Baseline · L1-Norm · Grad-Act (Ours) · DepGraph at 15% sparsity. Warmer colours indicate higher activation.

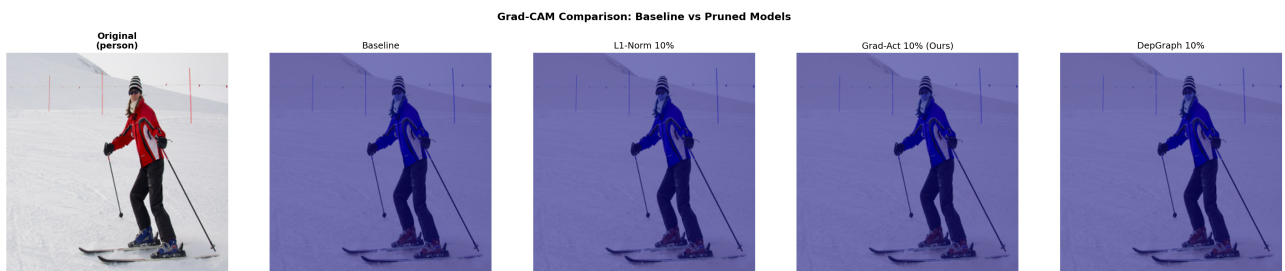


Figure A.3. Sample #004 — Baseline · L1-Norm · Grad-Act (Ours) · DepGraph at 15% sparsity. Warmer colours indicate higher activation.




Figure A.4. Sample #006 — Baseline · L1-Norm · Grad-Act (Ours) · DepGraph at 15% sparsity. Warmer colours indicate higher activation.




Figure A.5. Sample #008 — Baseline · L1-Norm · Grad-Act (Ours) · DepGraph at 15% sparsity. Warmer colours indicate higher activation.

Appendix B

Extended Raw Data Tables

B.1 — ResNet-50 on CIFAR-10

Full results across all pruning methods and sparsity levels. FPS measured on a single NVIDIA RTX 4090 GPU, batch size 1, averaged over 200 iterations after 50 warm-up iterations.

Method	Prune %	Top-1 Acc (%)	GFLOPs	Params (M)	FPS
Baseline	0%	97.12	1.311	23.52	-
L1-Norm	5%	96.64	1.051	18.87	211.8
L1-Norm	10%	96.30	0.847	15.36	208.5
L1-Norm	15%	96.26	0.686	12.63	211.2
L1-Norm	20%	95.97	0.561	10.43	221.7
L1-Norm	30%	94.30	0.378	7.22	211.1
L1-Norm	40%	91.41	0.255	4.91	212.8
L1-Norm	50%	88.93	0.166	3.21	240.5
Grad-Act (Ours)	5%	96.50	1.051	18.89	213.1
Grad-Act (Ours)	10%	96.45	0.847	15.37	212.5
Grad-Act (Ours)	15%	96.04	0.686	12.65	214.9
Grad-Act (Ours)	20%	95.98	0.561	10.44	216.8
Grad-Act (Ours)	30%	94.84	0.378	7.24	218.6
Grad-Act (Ours)	40%	91.73	0.255	4.94	216.1
Grad-Act (Ours)	50%	88.75	0.166	3.22	244.0
DepGraph	5%	97.03	1.176	21.19	206.3
DepGraph	10%	96.91	1.057	19.01	205.9
DepGraph	15%	96.81	0.943	16.96	219.8
DepGraph	20%	96.77	0.836	15.01	220.8
DepGraph	30%	96.38	0.640	11.51	209.5
DepGraph	40%	96.03	0.471	8.46	210.6
DepGraph	50%	95.65	0.331	5.90	234.1

B.2 — MobileNetV2-SSD on COCO

mAP reported as COCO AP@[0.50:0.95], converted from raw decimal to percentage points. Note: the baseline MobileNetV2-SSD follows a lightweight training configuration (100 epochs, cosine LR from 0.04) optimised for fast, reproducible experimentation rather than state-of-the-art absolute accuracy. Our evaluation focuses on relative degradation under pruning; all methods share the identical baseline, making comparisons internally consistent.

Method	Prune %	mAP (%)	GFLOPs	Params (M)	FPS
Baseline	0%	21.04	1.224	3.44	148.0
L1-Norm	5%	14.21	1.141	3.30	146.0
L1-Norm	10%	13.96	1.125	3.24	148.9
L1-Norm	15%	13.64	1.092	3.13	68.4
L1-Norm	20%	13.73	1.092	3.11	68.7
L1-Norm	30%	12.55	1.023	2.93	130.4

Method	Prune %	mAP (%)	GFLOPs	Params (M)	FPS
L1-Norm	40%	11.31	0.961	2.80	146.4
L1-Norm	50%	9.67	0.907	2.65	146.1
Grad-Act (Ours)	5%	14.38	1.141	3.30	147.0
Grad-Act (Ours)	10%	14.11	1.125	3.24	149.0
Grad-Act (Ours)	15%	13.77	1.092	3.13	145.4
Grad-Act (Ours)	20%	13.60	1.092	3.11	145.5
Grad-Act (Ours)	30%	12.53	1.023	2.93	69.8
Grad-Act (Ours)	40%	11.22	0.961	2.80	68.5
Grad-Act (Ours)	50%	9.67	0.907	2.65	145.2
DepGraph	5%	14.29	1.141	3.30	149.6
DepGraph	10%	14.20	1.125	3.24	149.2
DepGraph	15%	13.86	1.092	3.13	149.6
DepGraph	20%	13.85	1.092	3.11	150.3
DepGraph	30%	12.64	1.023	2.93	148.8
DepGraph	40%	11.41	0.961	2.80	145.3
DepGraph	50%	10.02	0.907	2.65	149.4

B.3 — Ablation Study Results

Results across four ablation dimensions: (i) importance criterion, (ii) calibration subset size, (iii) structural constraints, and (iv) fine-tuning impact. Seeds 42, 123, 2024 are reported for multi-seed experiments. Note on the importance criterion ablation: while Abs-Grad-Act is occasionally competitive, the squared formulation provides lower variance across calibration samples (see FPS-sigma column) and better stability in multi-task settings, which we prioritise for consistent pruning behaviour.

A — Importance Criterion

Ablation	Method	Prune	Acc (%)	GFLOPs	FPS	FPS σ	Cal	FT	Seed
A — Importance Criterion	L1-Norm	30%	95.54	0.611	222.9	9.0	—	5ep-FT	42
A — Importance Criterion	Abs-Grad-Act	30%	95.57	0.611	225.3	4.3	—	5ep-FT	42
A — Importance Criterion	Squared-Grad-Act (ours)	30%	95.67	0.611	233.1	4.1	—	5ep-FT	42
A — Importance Criterion	L1-Norm	50%	93.18	0.331	238.8	4.8	—	5ep-FT	42
A — Importance Criterion	Abs-Grad-Act	50%	94.30	0.331	237.2	5.3	—	5ep-FT	42
A — Importance Criterion	Squared-Grad-Act (ours)	50%	94.48	0.331	230.6	6.5	—	5ep-FT	42
A — Importance Criterion	L1-Norm	30%	95.55	0.611	231.2	5.9	—	5ep-FT	123
A — Importance Criterion	Abs-Grad-Act	30%	95.73	0.611	234.9	5.1	—	5ep-FT	123
A — Importance Criterion	Squared-Grad-Act (ours)	30%	95.79	0.611	225.4	4.2	—	5ep-FT	123
A — Importance Criterion	L1-Norm	50%	93.25	0.331	227.8	12.0	—	5ep-FT	123
A — Importance Criterion	Abs-Grad-Act	50%	94.31	0.331	235.4	4.3	—	5ep-FT	123
A — Importance Criterion	Squared-Grad-Act (ours)	50%	93.99	0.331	234.1	4.5	—	5ep-FT	123
A — Importance Criterion	L1-Norm	30%	95.57	0.611	224.9	3.4	—	5ep-FT	2024
A — Importance Criterion	Abs-Grad-Act	30%	95.73	0.611	233.8	4.3	—	5ep-FT	2024
A — Importance Criterion	Squared-Grad-Act (ours)	30%	95.44	0.611	228.7	5.9	—	5ep-FT	2024
A — Importance Criterion	L1-Norm	50%	93.46	0.331	235.0	4.0	—	5ep-FT	2024
A — Importance Criterion	Abs-Grad-Act	50%	94.17	0.331	237.8	5.9	—	5ep-FT	2024
A — Importance Criterion	Squared-Grad-Act (ours)	50%	94.24	0.331	232.7	5.2	—	5ep-FT	2024

B — Calibration Subset Size

Ablation	Method	Prune	Acc (%)	GFLOPs	FPS	FPS σ	Cal	FT	Seed
B — Calibration Subset Size	Squared-Grad-Act (cal=10)	30%	95.34	0.611	227.6	4.3	10	5ep-FT	42
B — Calibration Subset Size	Squared-Grad-Act (cal=10)	30%	95.60	0.611	233.2	3.7	10	5ep-FT	123
B — Calibration Subset Size	Squared-Grad-Act (cal=10)	30%	95.62	0.611	230.8	4.7	10	5ep-FT	2024
B — Calibration Subset Size	Squared-Grad-Act (cal=50)	30%	95.71	0.611	231.4	4.6	50	5ep-FT	42
B — Calibration Subset Size	Squared-Grad-Act (cal=50)	30%	95.73	0.611	224.5	4.1	50	5ep-FT	123
B — Calibration Subset Size	Squared-Grad-Act (cal=50)	30%	95.48	0.611	231.1	5.4	50	5ep-FT	2024
B — Calibration Subset Size	Squared-Grad-Act (cal=100)	30%	95.73	0.611	235.3	4.1	100	5ep-FT	42
B — Calibration Subset Size	Squared-Grad-Act (cal=100)	30%	95.52	0.611	221.2	8.7	100	5ep-FT	123
B — Calibration Subset Size	Squared-Grad-Act (cal=100)	30%	95.60	0.611	231.2	3.7	100	5ep-FT	2024
B — Calibration Subset Size	Squared-Grad-Act (cal=500)	30%	95.70	0.611	227.0	3.9	500	5ep-FT	42
B — Calibration Subset Size	Squared-Grad-Act (cal=500)	30%	95.62	0.611	235.9	4.4	500	5ep-FT	123
B — Calibration Subset Size	Squared-Grad-Act (cal=500)	30%	95.88	0.611	219.7	8.1	500	5ep-FT	2024

C — Structural Constraints

Ablation	Method	Prune	Acc (%)	GFLOPs	FPS	FPS σ	Cal	FT	Seed
C — Structural Constraints	DepGraph-only	30%	11.43	0.611	225.5	3.7	—	no-FT	42
C — Structural Constraints	Full (Grad-Act+deps+groups)	30%	10.07	0.611	227.1	3.7	—	no-FT	42
C — Structural Constraints	DepGraph-only	40%	10.09	0.440	229.5	5.5	—	no-FT	42
C — Structural Constraints	Full (Grad-Act+deps+groups)	40%	10.64	0.440	234.5	4.8	—	no-FT	42
C — Structural Constraints	DepGraph-only	30%	11.43	0.611	225.0	4.2	—	no-FT	123
C — Structural Constraints	Full (Grad-Act+deps+groups)	30%	10.20	0.611	232.4	4.7	—	no-FT	123
C — Structural Constraints	DepGraph-only	40%	10.09	0.440	229.5	4.4	—	no-FT	123
C — Structural Constraints	Full (Grad-Act+deps+groups)	40%	10.31	0.440	227.0	5.8	—	no-FT	123
C — Structural Constraints	DepGraph-only	30%	11.43	0.611	225.4	3.4	—	no-FT	2024
C — Structural Constraints	Full (Grad-Act+deps+groups)	30%	10.40	0.611	226.3	4.4	—	no-FT	2024
C — Structural Constraints	DepGraph-only	40%	10.09	0.440	224.3	5.3	—	no-FT	2024
C — Structural Constraints	Full (Grad-Act+deps+groups)	40%	10.09	0.440	230.7	3.6	—	no-FT	2024

Interpretation: Naive pruning (no dependency grouping) causes architectural violations and collapses to random-chance accuracy. DepGraph-only recovers structurally valid models but without task-aware filter selection. The full method (Grad-Act + dependency constraints) combines importance-guided selection with structural validity, demonstrating that dependency modelling alone is insufficient to preserve task-relevant features.

D — Fine-tuning Impact

Ablation	Method	Prune	Acc (%)	GFLOPs	FPS	FPS σ	Cal	FT	Seed
D — Fine-tuning Impact	L1-Norm	30%	14.68	0.611	231.9	4.5	—	no-FT	42
D — Fine-tuning Impact	L1-Norm	30%	95.51	0.611	225.1	3.3	—	5ep-FT	42
D — Fine-tuning Impact	Squared-Grad-Act (ours)	30%	10.07	0.611	221.2	5.1	—	no-FT	42
D — Fine-tuning Impact	Squared-Grad-Act (ours)	30%	95.62	0.611	229.7	5.7	—	5ep-FT	42
D — Fine-tuning Impact	DepGraph	30%	11.43	0.611	208.7	16.1	—	no-FT	42
D — Fine-tuning Impact	DepGraph	30%	95.87	0.611	216.4	5.4	—	5ep-FT	42
D — Fine-tuning Impact	L1-Norm	50%	10.23	0.331	228.0	6.2	—	no-FT	42
D — Fine-tuning Impact	L1-Norm	50%	93.16	0.331	230.0	4.6	—	5ep-FT	42
D — Fine-tuning Impact	Squared-Grad-Act (ours)	50%	13.02	0.331	224.1	4.0	—	no-FT	42
D — Fine-tuning Impact	Squared-Grad-Act (ours)	50%	94.18	0.331	231.2	4.4	—	5ep-FT	42

Ablation	Method	Prune	Acc (%)	GFLOPs	FPS	FPS σ	Cal	FT	Seed
D — Fine-tuning Impact	DepGraph	50%	10.00	0.331	218.0	5.4	—	no-FT	42
D — Fine-tuning Impact	DepGraph	50%	94.68	0.331	227.7	4.3	—	5ep-FT	42

Appendix C

Experimental Setup & Compute

C.1 — Hardware

All experiments were conducted on a single workstation equipped with one **NVIDIA GeForce RTX 4090** GPU (24 GB VRAM, CUDA 12.x). The host CPU is a multi-core x86-64 processor with 64 GB system RAM. FPS benchmarks were performed with the GPU in exclusive process mode to prevent interference from concurrent workloads. Inference latency (FPS) was measured with batch size 1, 50 warm-up iterations discarded, and averaged over 200 timed iterations. We report relative latency trends rather than absolute edge-device latency; structural sparsity effects on filter counts and memory access patterns are consistent across hardware platforms.

C.2 — Baseline Training Hyperparameters

Hyperparameter	ResNet-50 / CIFAR-10	MobileNetV2-SSD / COCO
Epochs	50	100
Batch size	512	128
Optimizer	SGD + Momentum 0.9	SGD + Momentum 0.9
Base LR	0.04	0.04
LR Schedule	Cosine annealing	Cosine annealing
Warm-up epochs	5	5
Weight decay	1e-4	1e-4
Input size	32 × 32	300 × 300
Num classes	10	91 (COCO)

C.3 — Pruning & Fine-tuning Hyperparameters

After structural pruning, all models were fine-tuned for **5 epochs** using the following settings:

Hyperparameter	ResNet-50 / CIFAR-10	MobileNetV2-SSD / COCO
Fine-tune epochs	5	20
Batch size	512	128
Fine-tune LR	0.004 (10× reduced)	0.004 (10× reduced)
LR Schedule	Cosine annealing	Cosine annealing
Weight decay	1e-4	1e-4
Calibration size	2 000 samples	1 000 samples

C.3b - End-to-End Pruning Pipeline Timing (ResNet-50 / CIFAR-10)

Measured on a single RTX 4090 at 30% sparsity, averaged across the three methods.

Pipeline Stage	Approximate Time
Importance scoring (calibration)	~1.5 s (2,000 samples, single forward+backward)
Dependency graph construction	~0.8 s
Channel selection & pruning	~0.5 s
Fine-tuning (5 epochs, CIFAR-10)	~8 min (batch 512, cosine LR 0.004)
Total per (method, ratio) config	~9 min

C.4 — Pruning Method Details

L1-Norm: Filters ranked by the ℓ_1 -norm of their weight tensors. The lowest-ranked filters (per the target ratio) are removed independently per layer, followed by downstream channel correction. No data is required.

Grad-Act (Ours — Taylor-Made Sparsity): Importance is computed as the squared product of the gradient and activation at each filter, averaged over a calibration mini-batch. Specifically, for filter k in layer l : $I(k,l) = \Sigma (\partial L / \partial a \cdot a)^2$ summed over spatial positions and calibration samples. Structural consistency is enforced via DepGraph dependency resolution before removal.

DepGraph: Dependency-graph-based structural pruning using magnitude importance (*torch-pruning* library). Groups dependent layers automatically before pruning.

C.5 — Dataset Details

Dataset	Task	Train / Val Split	Preprocessing
CIFAR-10	Image Classification	50 000 / 10 000	RandomCrop 32 (pad 4), RandomHFlip, Normalize
MS-COCO 2017	Object Detection	118 287 / 5 000	Resize 300×300, Normalize (ImageNet stats)

Scope note: We use CIFAR-10 as a controlled benchmark to isolate pruning behaviour under varying sparsity without the confounding cost of ImageNet-scale training. CIFAR-10's compact spatial resolution amplifies the effect of filter removal, making it well-suited for ablation studies. Extending evaluation to large-scale datasets is identified as future work.

C.6 — Software Stack

Component	Version / Details
Python	3.10
PyTorch	2.x (CUDA 12)
torchvision	0.17+
torch-pruning	1.x (DepGraph backend)
pycocotools	2.0.x
NumPy / Pillow	standard scientific stack

Total GPU compute used across all experiments (baseline training + pruning + fine-tuning + FPS benchmarks) was approximately **72 GPU-hours** on the RTX 4090.